



Text Analytics

Text Classification

Ulf Leser

Content of this Lecture

- Classification
- Algorithms
- Case studies

Disclaimer

- This is not a course on Foundations of **Machine Learning**
- Classification/clustering are presented rather briefly
 - There exist **many more methods**, much work on **comparing them** empirically, and a lot of work on **explaining the differences** between the different approaches
- General experience: Choosing another classification / clustering typically will not lead to dramatic improvements
 - Instances are either well classifiable or not
 - Changing the classification method **may yield 5-10% improvement**, but usually not more
- More important: Choice of features
 - This requires creativity and must be adapted to every problem
 - We do not discuss **feature selection**

Text Classification

- Given a set D of docs and a set of classes C . A **classifier** is a function $f: D \rightarrow C$
- How does this work in general?
 - Find a function v that maps a doc into a **vector of features**
 - For instance, its bag-of-words, possibly weighted by TF*IDF
 - Obtain a set D of docs with their classes
 - Find the characteristics of the features of docs in each class (= **build a model**)
 - What do they have in common?
 - How do they differ from docs in other classes?
 - Encode the model in a classifier function f operating on a feature vector: $v: D \rightarrow V$, and $f: V \rightarrow C$
 - We compute $f(v(d))$

Good Classifiers

- Our problem: Finding a **good classifier**
 - A good classifier assigns as many docs as possible to their “correct class”
- How do we know?
 - **Supervised learning**
 - Classification needs a **sample S** of docs with their correct classes
 - S is required for
 - Learning the model
 - Evaluating f : f is the better, the more docs are assign **their correct class**
 - Details on evaluation methods later

Overfitting

- We can easily build a **perfect classifier for S**
 - $f(d) = \{f(d'), \text{ if } \exists d' \in S \text{ with } d'=d; \text{ random otherwise}\}$
 - Applied to only docs from S, f is a perfect classifier
- But: This classifier will not work well on “new” documents
- Improvement
 - $f(d) = \{f(d'), \text{ if } \exists d' \in S \text{ with } d' \sim d; \text{ random otherwise}\}$
 - If S is small and “ \sim ” very narrow, this does not help a lot
 - But see kNN classifiers
- **Overfitting**
 - If the **model strongly depends on S**, f overfits – it will only work well if all future d’s are very similar to the docs in S
 - You cannot find overfitting when **evaluation is performed on S** only

Against Overfitting

- **f must generalize**: Capture features that are typical for all docs in D , not for the docs in S
- Still, we only have S for evaluation ...
 - We need to extrapolate the quality of f to **unknown docs**
- Usual method: **Cross-validation** (leave-one-out, jack-knife)
 - Partition S into k sets (typical: $k=10$)
 - Leave-one-out: $k=|S|$
 - Learn model on $k-1$ sets and evaluate on the k 'th
 - Perform k times, each time evaluating on another partition
 - Estimated quality on new docs = **average performance**
 - Often the best we can do

Problem 1: Information Leakage

- Developing a classifier is an **iterative process**
 - Define feature vector
 - Evaluate performance using cross-validation
 - Perform error analysis, leading to others features
 - Iterate until satisfied with result
- In this process, you “sneak” into the data (during error analysis) you later will evaluate on
 - “**Information leakage**”: Information on eval data is used in training
- Solution
 - Reserve a portion P of S for evaluation
 - Perform iterative process only on $S \setminus P$
 - Final evaluation on P ; **then no more iterations**

Problem 2: Biased S

- Very often, S is biased
 - Often, one class c' (or some classes) is **much less frequent** than the other(s)
 - E.g. finding text written in dialect
 - To have enough inst. of c' in S , these are searched actively in D
 - Later, examples from other classes are added
 - But how many?
 - **Fraction of c' in S** is much (?) higher than expected by chance
 - I.e., than obtained by random sampling
- Solutions
 - Try to estimate **fraction of c' in D** and produce stratified S
 - Very difficult and costly, often almost impossible
 - Because S would need to be very large


A Simple Example

- An aggregated history of credit loss in a bank

Class	Age	Income	Risk
1	20	1500	High
2	30	2000	Low
3	35	1500	High
4	40	2800	Low
5	50	3000	Low
6	60	6000	High

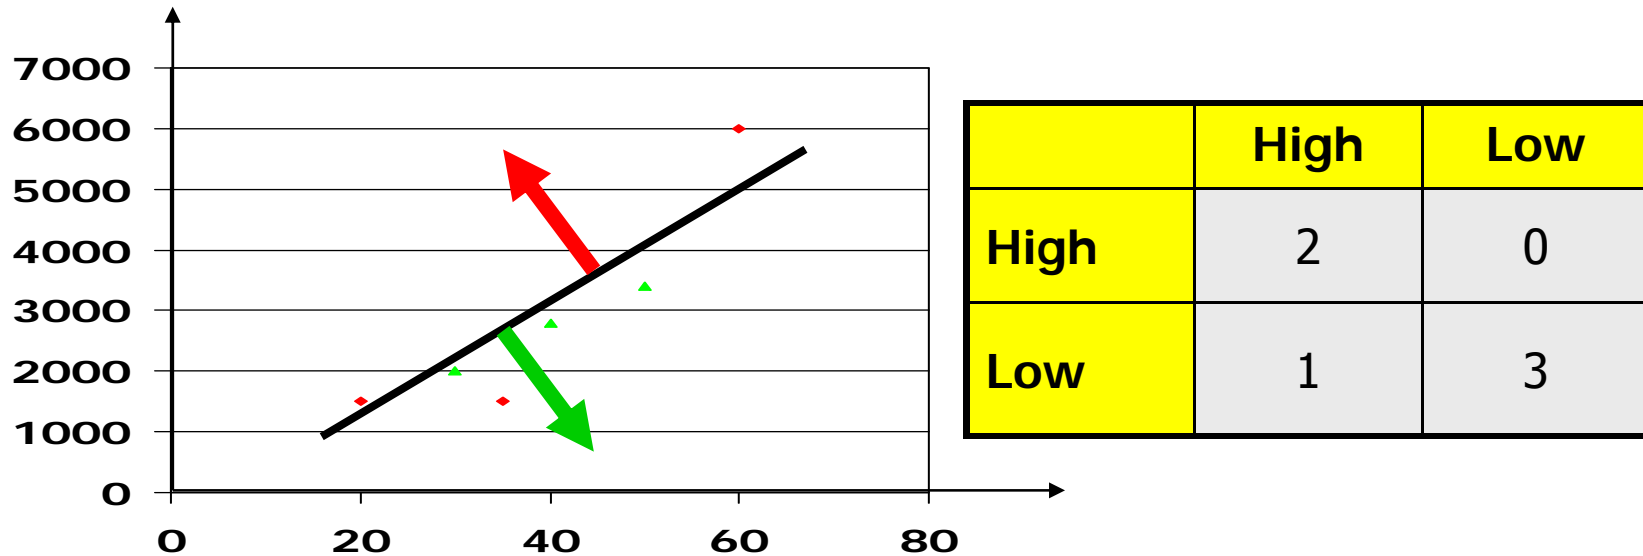
- Now we see a **new person**, 45 years old, 4000 Euro income
- What is his risk?

Regression



- Simple approach: **Linear regression**
 - Linear separation with minimum square of error
- Use location relative to **regression line as classifier**
- Compute parameters such that error is the smallest
 - This is one way of doing it; no details on regression here

Performance on the Training Data



- Quality of predicting “high risk”
 - Precision = $TP/(TP+FP) = 2/2$, Recall = $TP/(TP+FN) = 2/3$, Accuracy: $5/6$
- Regression makes **many assumptions**
 - Assumes linear correlations between attributes
 - Requires **numerical attributes**
 - Method of choice if C is continuous (infinitely many ordered classes)

Categorical Attributes

Class	Age	Type of car	Risk of Accident
1	23	Family	High
2	17	Sports	High
3	43	Sports	High
4	68	Family	Low
5	25	Truck	Low

- Assume this classification was created by some insurance manager. What was in his head?
 - Probably a **set of rules**, such as

```
if      age > 50      then risk = low
elseif age < 25      then risk = high
elseif car = sports  then risk = high
else    risk = low
```

Decision Rules

Class	Age	Type of car	Risk of Accident
1	23	Family	High
2	17	Sports	High
3	43	Sports	High
4	68	Family	Low
5	25	Truck	Low

- Can we find **less rules** which, for these data sets, result in the same classification?

```
if      age > 50      then risk = low
elseif car = truck then risk = low
else    risk = high
```

A Third Approach

Class	Age	Type of car	Risk of Accident
1	23	Family	High
2	17	Sports	High
3	43	Sports	High
4	68	Family	Low
5	25	Truck	Low

- Why not:

```
If      age=23 and car = family then risk = high
elseif age=17 and car = sports then risk = high
elseif age=43 and car = sports then risk = high
elseif age=68 and car = family then risk = low
elseif age=25 and car = truck  then risk = low
else    flip a coin
```

Overfitting - Again

- This was in instance of our “perfect classifier”
- We always learn a model from a **small sample** of the real world
- **Overfitting**
 - If the model is too close to the training data, it performs perfect on the training data but learned any bias present in the training data
 - Thus, the rules **do not generalize** well
- **Solution**
 - Use an appropriate feature set and learning algorithm
 - Evaluate you method using cross-validation

Text Classification

- Many problems in text analytics can be cast as classification
 - Language identification
 - Topic identification
 - Spam detection
 - Content-based message routing
 - Named entity recognition (is this token part of a NE?)
 - Author identification (which plays were really written by Shakespeare?)
 - ...
- Common problem
 - No well discriminating single features
 - We need to use a high dimensional feature space

Classification Methods

- There are a **zillion different methods**
 - k-nearest neighbor
 - Naïve Bayes and Bayesian Networks
 - Decision Trees and Rainforests
 - Maximum Entropy, Maximum Entropy Markov Models, Conditional Random Fields
 - Support Vector Machines
 - Perceptrons, Neural Networks
 - ...
- **Effectiveness of classification** depends on problem, algorithm, feature selection method, sample, evaluation, ...
 - But: Often the difference between different methods are astonishing small

Content of this Lecture

- Classification
- Algorithms
 - Nearest Neighbor
 - Naïve Bayes
 - Maximum Entropy
 - Linear Models and Support Vector Machines (SVM)
- Case studies

Nearest Neighbor Classifiers

- Very simple and **effective method**

- Definition


Let S be a set of classified documents, m a distance function between any two documents, and d an unclassified doc.

- A **nearest-neighbor (NN) classifier** assigns to d the class of the nearest document in S (wrt. m)
- A **k -nearest-neighbor (kNN) classifier** assigns to d the most frequent class among the k nearest documents in $(S \text{ wrt. } m)$


- Remark

- Obviously, a proper **distance function** is very important
- We may weight the k nearest docs according to their distance to d
- We need to take care of multiple docs with the same distance

Illustration



1NN and **Voronoi diagram** marking the regions of influence around each data point (2D-space, each point a separate class)



A 5NN

Properties

- Assumption: **Similar docs should have the same class**
 - Depends a lot on the distance function
- kNN is simple and **astonishing good**
- kNN in general is more robust than NN
- (k)NN is an example of **lazy learning**
 - Actually, there is no learning
 - Actually, there is no model
 - Where are the features?
- Features
 - We still need to define features
 - These features are the **input to the distance function**

Disadvantages

- Major problem: **Performance** (speed)
 - We need to compute the distance between d and any doc in S
 - This requires $d \cdot |S|$ applications of the distance function
 - Often the cosine of two 100K-dimensional vectors
- Various suggestions for speeding-up the method
 - Clustering – aggregate groups of **very close points in S** into a single representative
 - Linear speed-up
 - Extreme case: Chose one representative per class
 - Usually not a good idea (high dimensional space!); no kNN any more; very fast and space efficient
 - Multidimensional index structures and **metric embeddings**
 - Map into a lower-dimensional space such that distances are preserved

kNN for Text

- In the VSM world, kNN is implemented very easily using the methods we already learned
- How?
 - Use cosine distance of bag-of-word vectors as distance
 - The usual VSM query mechanism computes exactly the k nearest neighbors when d is used as query
 - Difference
 - d usually much larger than the average q
 - We might need other ways of optimizing “queries”

Content of this Lecture

- Classification
- Algorithms
 - Nearest Neighbor
 - Naïve Bayes
 - Maximum Entropy
 - Linear Models and Support Vector Machines (SVM)
- Case studies

Bayes' Classification

- Simple method based on **relative frequencies of features in the different classes**
- Given
 - Set S of docs and set of classes $C = \{c_1, c_2, \dots, c_m\}$
 - Docs are described as a set F of **binary features**
 - Usually the presence/absence of terms in d
- We seek $p(c_i|d)$, the probability of a doc $d \in S$ being a member of class c_i
- d eventually is assigned to c with $p(c|d) = \operatorname{argmax} p(c_i|d)$
- Replace d with feature representation

$$p(c | d) = p(c | F[d]) = p(c | f_1[d], \dots, f_n[d]) = p(c | t_1, \dots, t_n)$$

Probabilities

- What we learn from the training data (MLE)
 - The **a-priori probability** $p(t)$ of every term t
 - How many docs from S have t ?
 - The **a-priori probability** $p(c)$ of every class $c \in C$
 - How many docs in S are of class c ?
 - The **conditional probabilities** $p(t|c)$ for term t being true in class c
 - Proportion of docs in c with term t among all docs in c
- Rephrase and use Bayes' theorem

$$p(c | t_1, \dots, t_n) = \frac{p(t_1, \dots, t_n | c) * p(c)}{p(t_1, \dots, t_n)} \approx p(t_1, \dots, t_n | c) * p(c)$$



Term can be dropped; value is **identical for all classes**, and we only want to rank the $p(c|d)$

Naïve Bayes

- We have $p(c | d) \approx p(t_1, \dots, t_n | c) * p(c)$
- The first term cannot be learned with any reasonably large training set
 - There are 2^n combinations of feature values
- Solution: Be „naïve“
 - Assume **statistical independence** of all terms
- Then $p(t_1, \dots, t_n | c) = p(t_1 | c) * \dots * p(t_n | c)$
- And finally

$$p(c | d) \approx p(c) * \prod_{i=1}^n p(t_i | c)$$

Properties

- Simple algorithm, quite robust, comparably fast, needs **extensive smoothing**
- Often used as **baseline** for other methods
- Learning the model is simple, and the model is quite compact ($O(|K|*|C|)$ space)
- When we use the logarithm (equally well for ranking), we see that NB is a **(log-)linear classifier**

$$\begin{aligned} p(c | d) &\approx \log\left(p(c) * \prod p(t_i | c)\right) \\ &= \log(p(c)) + \sum \log(p(t_i | c)) \end{aligned}$$

Feature Selection

- One can easily speed-up classification by using only a **subset of all features**
- Simplest case: Use those t where $p(t|c)$ show the **biggest differences** between the different classes
- Numerous methods for **feature selection**
 - Information gain, statistical tests, Bayesian information criterion, GINI score, ...
 - Finding the best features is not the same as finding the **best subset of features**
 - Overfitting is an issue: “Best features for S ” \neq “best features for D ”
- Same methods benefit from feature selection, some not
 - SVM usually not, Bayes usually yes (think of redundant features)

Content of this Lecture

- Classification
- Algorithms
 - Nearest Neighbor
 - Naïve Bayes
 - Maximum Entropy
 - Support Vector Machines (SVM)
- Case studies

Discriminative versus Generative Models

- NB uses Bayes' Theorem to estimate $p(c|d)$

$$p(c | t_1, \dots, t_n) = \frac{p(t_1, \dots, t_n | c) * p(c)}{p(t_1, \dots, t_n)} \approx p(t_1, \dots, t_n | c) * p(c)$$

- Notation

- Approaches that **estimate $p(d|c)$ are called generative**
 - $p(d|c)$ is the probability of class c producing data d
 - Thus, NB is a generative model
- Approaches that **directly estimate $p(c|d)$ are called discriminative**

Maximum Entropy Modeling

- Maximum Entropy (ME) is discriminative
- Given a set of **binary features**, it directly learns conditional probabilities $p(c|d)$
- Definition
*Let s_{ij} be the score of feature i for doc d_j (such as $TF*IDF$). We derive from s_{ij} a **binary indicator function** f_i for doc j and class c :*

$$f_i(d_j, c) = \begin{cases} 1, & \text{if } s_{ij} > 0 \wedge c = 1 \\ 0 & \text{otherwise} \end{cases}$$

- Remark
 - We will often call those indicator functions “features”, although they embed information about classes (“a **feature in a class**”)

Classification with ME

- Since $p(c,d)=p(c|d)*p(d)$ and $p(d)$ is the same for all c , we directly use $p(c|d)\sim p(c,d)$
- The ME approach models the **joint probability $p(c,d)$** as

$$p(c, d) = \frac{1}{Z} * \prod_{i=1}^K \alpha_i^{f_i(d, c)}$$

- Z is a normalization constant
 - The **feature weights α_i** are learned from the data
 - K is the number of features
- Classification with ME
 - Compute $p(c,d)$ for all c and return the **class with the highest value**

Finding Feature Weights

- Of course, the problem is finding appropriate α_i
- We want to choose the α_i such that the probability of the training data S given the model M is maximized

$$p(S | M) = \sum_{d \in S} p(c(d), d | M)$$

- This choice must take the dependencies between the features in the model into account
- Naïve Bayes computes α -like values independently for each feature and uses their linear combination for classification
 - This only works if statistical independence holds
 - For instance, using the same feature multiple times does bias the NB result

Maximum Entropy Models

- Essentially, ME applies a search strategy to find those α_i
- Problem: There are **indefinitely many combinations** of weights that may all give rise to the same maximal probability of S
- ME chooses the model with the **largest entropy**
 - Abstract formulation: The training data leaves too much freedom. We want to choose M such that all “undetermined” probability mass is distributed equally
 - ME tried to make **as few assumptions** as possible given the data
 - Such a distribution exists and is unique
 - The search strategy needs to take this into account

Entropy of a Distribution

- Let F be the feature space and M be an assignment of probabilities to each state in F . The **entropy of the probability distribution M** is defined as:

$$h(M) = - \sum_{s \in F} p(s | M) * \log(p(s | M))$$

- Thus, ME searches M such that
 - $P(S|M)$ is maximal and
 - $h(M)$ is maximal

Example [NLTK, see <http://nltk.googlecode.com/svn/trunk/doc/book/ch06.html>]

- Assume we have 10 different classes A-J and **no further knowledge**. Now we want to classify an document d. Which probabilities would you assign to the classes?

	A	B	C	D	E	F	G	H	I	J
(i)	10%	10%	10%	10%	10%	10%	10%	10%	10%	10%
(ii)	5%	15%	0%	30%	0%	8%	12%	0%	6%	24%
(iii)	0%	100%	0%	0%	0%	0%	0%	0%	0%	0%

- Model (i) **does not model more than we know**
- Model (i) also has maximal entropy

Example continued

- Now we learn that A is true in 55% of all cases. Which model do you chose?

	A	B	C	D	E	F	G	H	I	J
(iv)	55%	45%	0%	0%	0%	0%	0%	0%	0%	0%
(v)	55%	5%	5%	5%	5%	5%	5%	5%	5%	5%
(vi)	55%	3%	1%	2%	9%	5%	0%	25%	0%	0%

- Model (v) also has maximal entropy

Example continued

- We additionally learn that if the word “up” appears in a document, then there is an 80% chance that A or C are true. Furthermore, “up” is contained in 10% of the docs.
- This would result in the following model
 - We now **introduce features**
 - The 55% a-priori chance for A still holds

	A	B	C	D	E	F	G	H	I	J
+up	5.1%	0.25%	2.9%	0.25%	0.25%	0.25%	0.25%	0.25%	0.25%	0.25%
-up	49.9%	4.46%	4.46%	4.46%	4.46%	4.46%	4.46%	4.46%	4.46%	4.46%

- Things get **more complicated** if we have >100k features

Example 2 [Pix,Stockschläder, WS07/08]

- Assume we count features “has blue eyes” and “is left-handed” among a population of tamarins
- We observe $p(\text{eye})=1/3$ and $p(\text{left})=1/3$
- What is the **joint probability** $p(\text{eye}, \text{blue})$ of blue-eyed, left-handed tamarins?
 - We don’t now
 - It must be $0 \leq p(\text{eye}, \text{blue}) \leq \min(p(\text{eye}), p(\text{left})) = 1/3$
- Four cases



$p(\dots, \dots)$	left-handed	not left-handed	sum
blue-eyed	x	$1/3 - x$	$1/3$
not blue-eyed	$1/3 - x$	$1 - 2/3 + x$	$2/3$
sum	$1/3$	$2/3$	1

Maximizing Entropy

- The **entropy of the joint distribution** M here is

$$h(M) = - \sum_{i=1}^4 p(x, y) * \log(p(x, y))$$

- The value is maximal for **$dH/dx = 0$**
- Computing the first derivative and solving the equation leads to $x=1/9$
 - Which, in this case, is the same as assuming independence, but this is not generally the case
 - In general, finding a solution in this analytical way is not possible

Generalized Iterative Scaling (idea only)

- How do we find M in general?
- Generalized Iterative Scaling
 - Iterative procedure finding the optimal solution
 - Essentially, it starts from a random guess of all the $p(c,d)$ and iteratively redistributes probability mass until convergence
 - See [MS99] for the algorithm
- Problem: Usually converges very slowly
 - Long training times
- Several improved algorithms are known
 - Improved Iterative Scaling
 - Conjugate Gradient Descent

Properties of Maximum Entropy Classifiers

- In general, ME outperforms NB
- ME **does not assume independence** of features
 - Feature weights are learned by always taking the **entire distribution** into account
 - Two “redundant” features will simply get half of the weight as if there was only one feature
- Very popular in statistical NLP
 - Some of the best POS-tagger are ME-based
 - Some of the best NER systems are ME-based
- Several extensions
 - Maximum Entropy Markov Models
 - Conditional Random Fields

Content of this Lecture


- Classification
- Algorithms
 - Nearest Neighbor
 - Naïve Bayes
 - Maximum Entropy
 - Support Vector Machines (SVM)
- Case studies

Class of Linear Classifiers

- Many common classifiers are (log-)linear classifiers
 - Naïve Bayes
 - Perceptron / Winnow
 - Linear and Logistic Regression
 - Maximum Entropy
 - Support Vector Machines
- If applied on a binary classification problem, all these methods somehow compute a hyperplane which (hopefully) separates the two classes
- Despite similarity, noticeable performance differences exist
 - Which of the infinite number of possible separating hyperplanes is chosen?
 - How are non-separable data sets (by a linear model) handled?
- Experience: Classifiers more powerful than linear often don't perform better (on text)

NB and Regression

- Using **linear regression**, we compute a separating hyperplane using error minimization
- If we assume **binary Naïve Bayes**, we may compute



$$\frac{p(c | d)}{p(\neg c | d)} \approx \log\left(\frac{p(c)}{p(\neg c)}\right) + \sum \log\left(\frac{p(t_i | c)}{p(t_i | \neg c)}\right)$$
$$= a + \sum_{k \in K} b_k * TF_k$$


This is a **linear hyperplane**; value > 0 gives c, value < 0 gives not c

ME is a Log-Linear Model

$$p(c, d) = \frac{1}{Z} * \prod_{i=1}^K \alpha_i^{f_i(d, c)} \approx \log\left(\frac{1}{Z}\right) + \sum_{i=1}^K f_i(d, c) * \alpha_i$$

Rocchio Classification

- Recall **relevance feedback** in the VSM using Rocchio
 - Compute initial result
 - **Build new query** by aggregating all true positives and discounting all/some false positives
- This idea can be turned into a classifier
 - Compute the **centroid of all positive examples**
 - Compute the centroid of all negative examples
 - Compute the **hyperplane with minimal distance** to both centroids




Text = High Dimensional Data

- Document co-ordinates are zero along almost all axes
- Most document pairs are **very far apart** (i.e., not strictly orthogonal, but only share very common words)
- In classification terms: **virtually all document sets are separable for essentially any classification**
 - This is part of why linear classifiers are quite successful in this domain
- The trick is more of finding the “right” separating hyperplane instead of just finding (any) one

Linear Classifiers


- **Hyperplane** separating classes in high dimensional space
 - For illustration, we stay in 2-dimensional
- But which?




Quelle: Xiaojin Zhu, SVM-cs540

Support Vector Machines (sketch only)

- SVMs compute the hyperplane which **maximizes the margin**
 - I.e., is as far away from any data point as possible
- Can be cast in a linear optimization problem and solved efficiently
 - Classification finally only depends on the **support vectors – efficient**
 - Points most closest to hyperplane
 - Complication since usually the **classes are not linearly separable**
 - Minimizes the error under some assumptions



Problems not Linearly Separable



- Map data into an even **higher dimensional space**
- Not-linearly separable sets may become linearly separable
- Doing this efficiently requires a good deal of work
 - The “kernel trick”

Properties of SVM

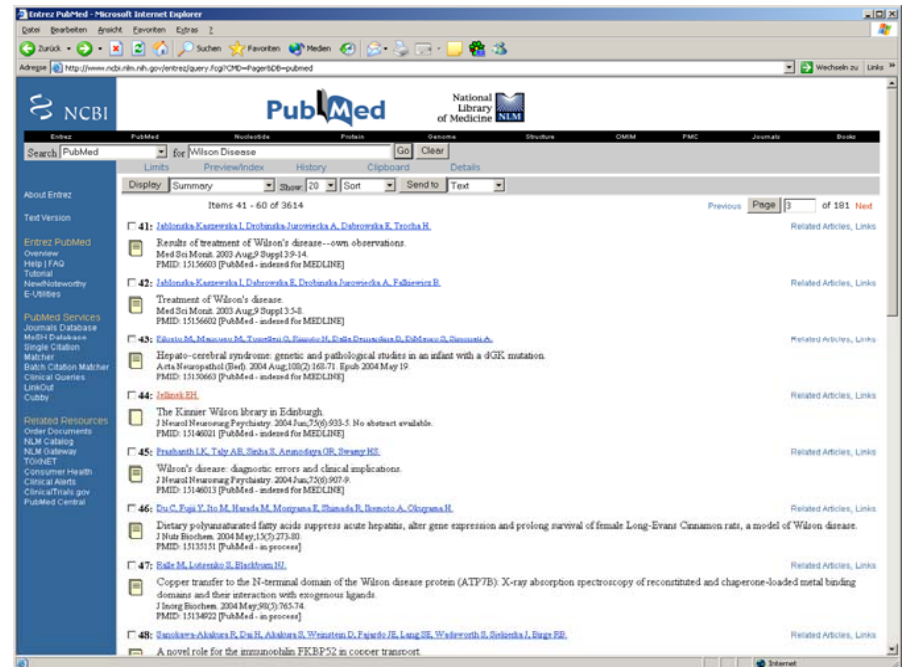
- State of the art in text classification
- Might require long training time
 - Worst case quadratic in training data
 - Various clever tricks and heuristics exist
- Classification is rather fast
 - Only distance to hyperplane is needed
 - Hyperplane is defined by only few vectors (support vectors)
- SVM are quite good “as is”, but lot of tuning possible
 - Kernel function, biased margins, ...
- Several implementations exist
 - SVMlight, libSVM, ...

Content of this Lecture

- Classification
- Algorithms
- Case studies
 - Topic classification
 - Spam filtering

Topic Classification [Rutsch et al., 2005]

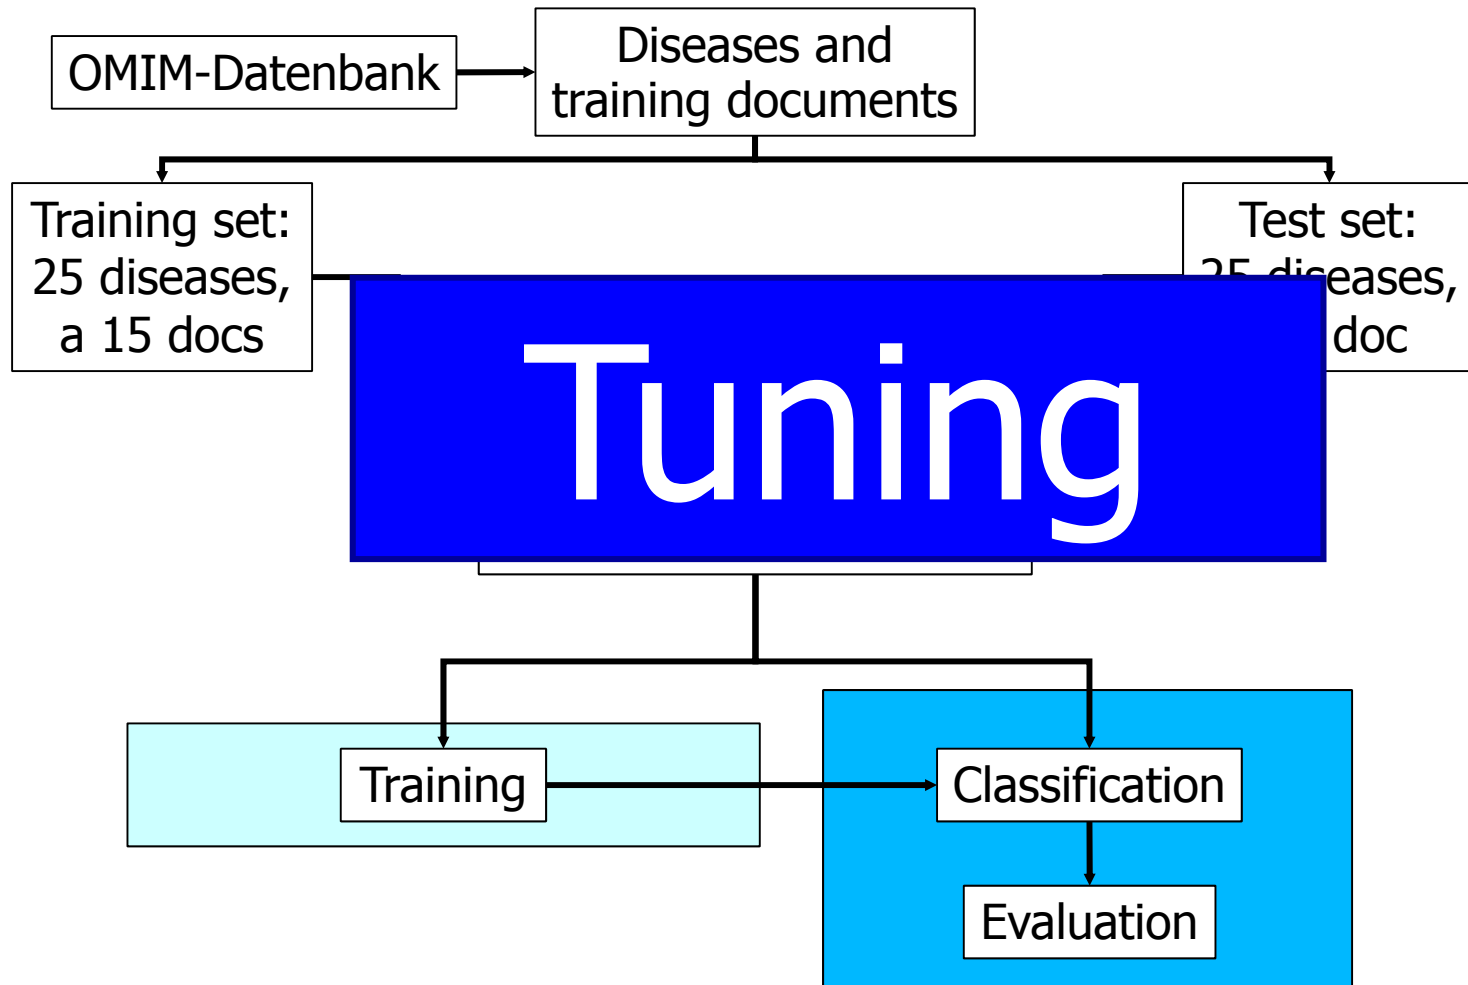
- Find publications treating the **molecular basis of hereditary diseases**
- Pure key word search generates too many results
 - "Asthma": 84 884 hits
 - Asthma and cats, factors inducing asthma, treatment, ...
 - "Wilson disease": 4552 hits
 - Including all publications from doctors named Wilson
- Pure key word search does not cope with **synonyms**




Idea

- Learn what is **typical for a paper** treating molecular basis of diseases from examples
 - 25 hereditary diseases
 - 20 abstracts for each disease
- We call this “typical” a **model** of the data
- Models are learned using some method
- Classification: Given a new text, find the model which fits best and **predict the associated class** (disease)
- What can we learn from 20 documents?

Complete Workflow




Results (Nearest-Centroid Classifier)



- Configurations (y-axis)
 - Stemming: yes/no
 - Stop words: 0, 100, 1000, 10000
 - Different forms of tokenization
- Best: No stemming, 10.000 stop words

Results with Section Weighting



- For fixed configuration, use **different weights** for terms depending on the section they appear in
 - Introduction, results, material and methods, discussion, ...

Influence of Stemming

Mit stemmer				
Nomen und Verben				
	100	1000	10000	
Precision	61,00	63,07	67,42	
Recall	59,29	60,51	65,01	
F-Measure	60,13	61,76	66,19	

Ohne Stemmer				
Nomen und Verben				
	100	1000	10000	
Precision	62,90	64,94	66,17	
Recall	62,59	62,38	62,71	
F-Measure	62,75	63,63	64,39	

Naive Bayes

- Best results

Versuche:	A	B	C	D	E
Precision	64.55	64.80	66.00	69.94	64.55
Recall	62.82	62.61	65.35	55.20	62.82
F-Measure	63.67	63.69	65.68	61.70	63.67

- A = Stemmer: ON, Stoppwörter: 10 000, Nomen-Tagging: ON, VerbenTagging: ON
- B = Stemmer: OFF, Stoppwörter : 10 000, Nomen-Tagging: ON, VerbenTagging: ON
- C = Stemmer: ON, Stoppwörter : 10 000, Nomen-Tagging: ON, VerbenTagging: OFF
- D = Stemmer: ON, Stoppwörter : OFF, Nomen-Tagging: OFF, VerbenTagging: OFF
- E = Stemmer: ON, Stoppwörter : 10 000, Nomen-Tagging: ON, VerbenTagging: ON

- **Nearest Centroid** outperforms Naïve Bayes
 - On this particular problem and training set ...


Content of this Lecture

- Classification
- Algorithms
- Case studies
 - Topic classification
 - Spam filtering

Thanks to: Conrad Plake, "Vi@gra and Co.: Approaches to E-Mail Spam Detection", Desden, December 2010

Spam

- Unsolicited Bulk E-Mail
- Old „problem“: 1978 first spam e-mail for advertisement
- Estimate: >95% of all mails are spam
- Many important issues not covered here
 - Filtering at provider, botnets, DNS filtering with black / gray / white lists, using further metadata (attachments, language, embedded images, n# of addressees, ...) etc.
 - Legal issues

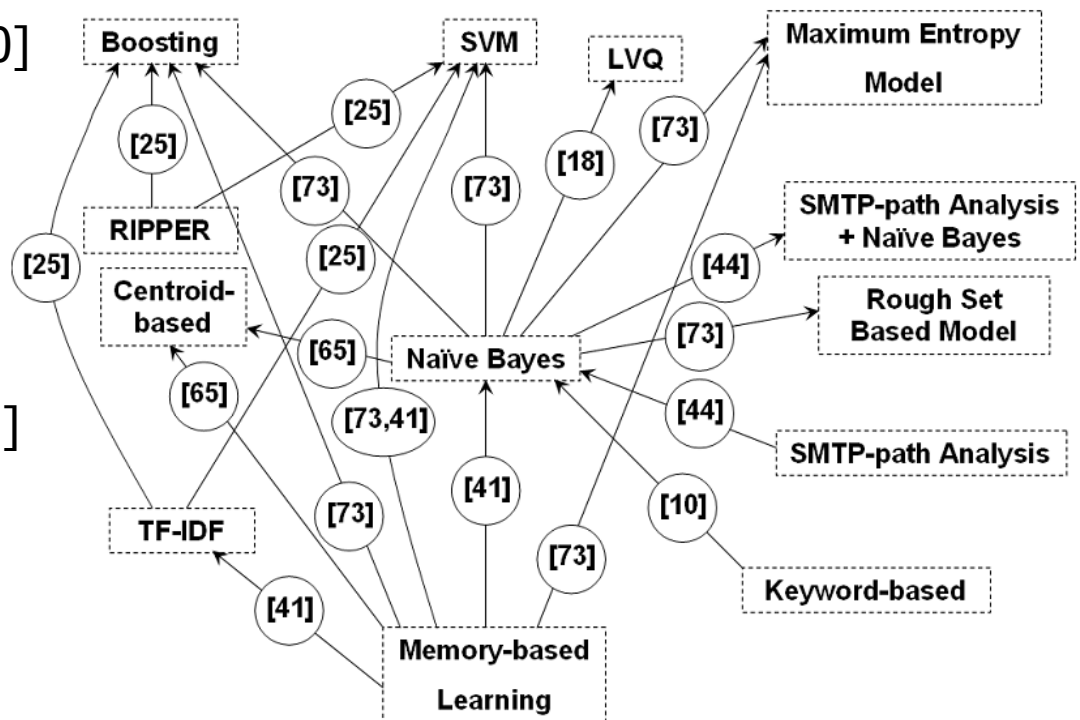


SPAM Detection as a Classification Task

- Content-based SPAM filtering
- Task: Given the body of an email – classify as SPAM or not
- Difficulties
 - Highly unbalanced classes (97% Spam)
 - Spammer react on every new trick – an arms race
 - Topics change over time
- Baseline approach: Naïve Bayes on VSM
 - Implemented in Thunderbird and MS-Outlook
 - Fast learning, relatively fast classification
 - Using TF, TF-IDF, Information Gain, ...
 - Stemming (mixed reports)
 - Stop-Word removal (seems to help)

Many Further Suggestions


- Rule learning
[Cohen, 1996]
- k-Nearest-Neighbors
[Androutsopoulos *et al.*, 2000]
- SVM
[Kolcz/Alspector, 2001]
- Decision trees
[Carreras/Marquez, 2001]
- Centroid-based
[Soonthornphisaj *et al.*, 2002]
- Artificial Neural Networks
[Clark *et al.*, 2003]
- Logistic regression
[Goodman/Yih, 2006]
- Maximum Entropy Models
- ...



Source: Blanzieri and Bryl, 2009

Measuring Performance

- We so far always assumed that a FP is as bad as a FN
 - Inherent in F-measure
- Is this true for Spam?
 - Missing a non-spam mail (FP) usually is perceived as much more severe than accidentally reading a spam mail (FN)
- Performance with growing feature sets and $c(\text{FP})=9*c(\text{FN})$



Problem Solved?

- Tricking a Spam filter
 - False feedback by malicious users (for global filters)
 - Bayesian attack: add "good" words
 - Change **orthography** (e.g., *viaagra*, *vi@gra*)
 - Tokenization attack (e.g., *free -> f r e e*)
 - **Image spam** (already >30%)
- Spam \neq Spam: **Concept drifts**
 - Spam topics change over time
 - Filters need to adapt



CEAS 2008 Challenge: Active Learning Task

- CEAS: Conference on Email and Anti-Spam
- Active Learning
 - Systems selected up to 1000 mails
 - Selection using score with **pre-learned model**
 - Classes of these were given
 - Simulates a system **which asks a user if uncertain**
- 143,000 mails

Name	Spam Caught %	Blocked Ham %	1-AUC %
Logistic Regression + Active Learning	99.92	0.12	0.0033
Online SVM (TREC07-tftS) - Entry 1	98.65	0.08	0.0250
Online SVM (TREC07-tftS) - Entry 3	98.65	0.07	0.0257
Heilongjiang Institute of Technology - Entry 3	98.66	0.14	0.0303
Online SVM (TREC07-tftS) - Entry 2	98.61	0.07	0.0331
Heilongjiang Institute of Technology - Entry 2	98.64	0.19	0.0557
PPM Compression (TREC07-ijspmm)	94.28	0.01	0.1031
Communication and Computer Network Lab (South China Univ. of Technology) - Entry 3	99.98	27.55	0.1500
Dynamic Markov Compression(TREC07-wat2)	98.11	0.34	0.2988
Communication and Computer Network Lab (South China Univ. of Technology) - Entry 2	99.88	25.53	0.5234
IGF (Ígor Assis Braga) - Entry 3	72.57	0.01	1.4495
IGF (Ígor Assis Braga) - Entry 2	80.59	0.01	8.9047
Kosmopoulos Aris - Entry 2	81.84	51.14	27.1210
Kosmopoulos Aris - Entry 1	86.20	57.20	28.7998